

# Clustering

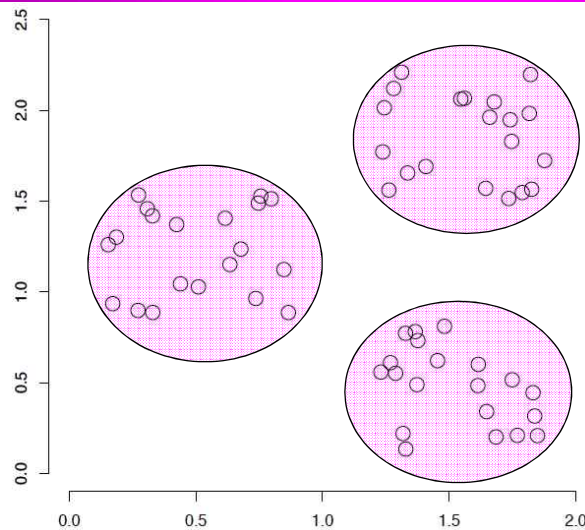
1

## What is clustering?

- **Clustering** is the process of grouping a set of objects into classes of similar objects
  - » Documents within a cluster should be similar.
  - » Documents from different clusters should be dissimilar.
- **Clustering** is the commonest form of *unsupervised learning*
  - » Unsupervised learning = learning from raw data where any classification of examples is not given
  - » Supervised learning = learning from supervised data where a classification of examples is given
- **Clustering** is a common and important task that finds many applications in IR and other places

2

## A data set with clear cluster structure



- How would you design an algorithm for finding these three clusters in this case?
- We are going to learn such clustering algorithms in this lecture!

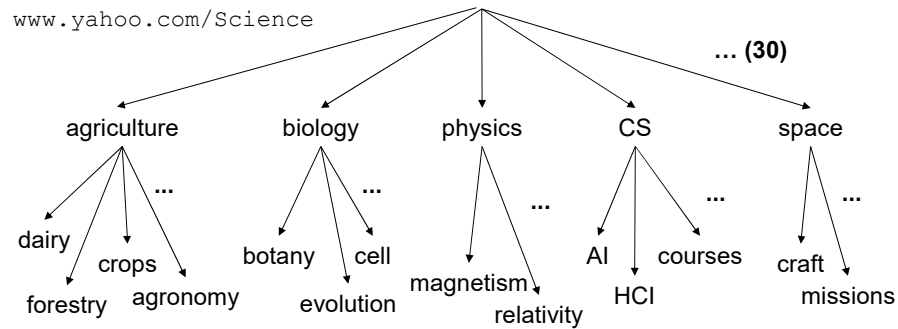
3

Clustering is widely used in many IR applications

4

## Topic Hierarchy of Yahoo!

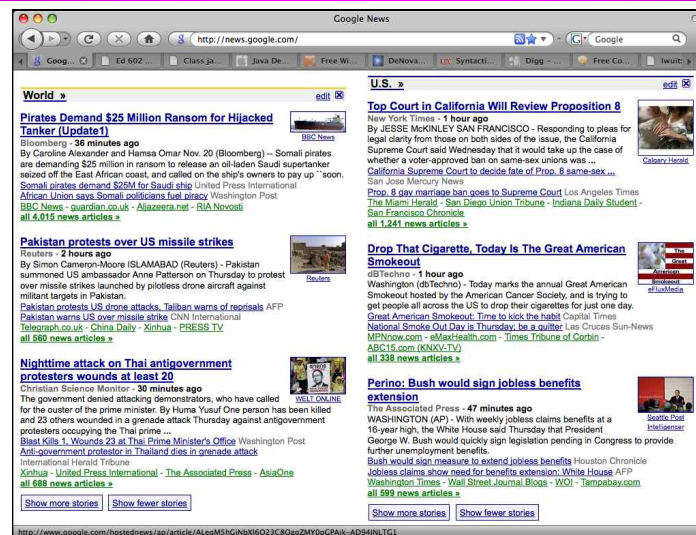
: the kind of output we obtain from clustering



Yahoo Science hierarchy, consisting of web pages that were gathered in summer of 1997 and pointed to by Yahoo  
This hierarchy has totally 264 classes, but only samples are shown here

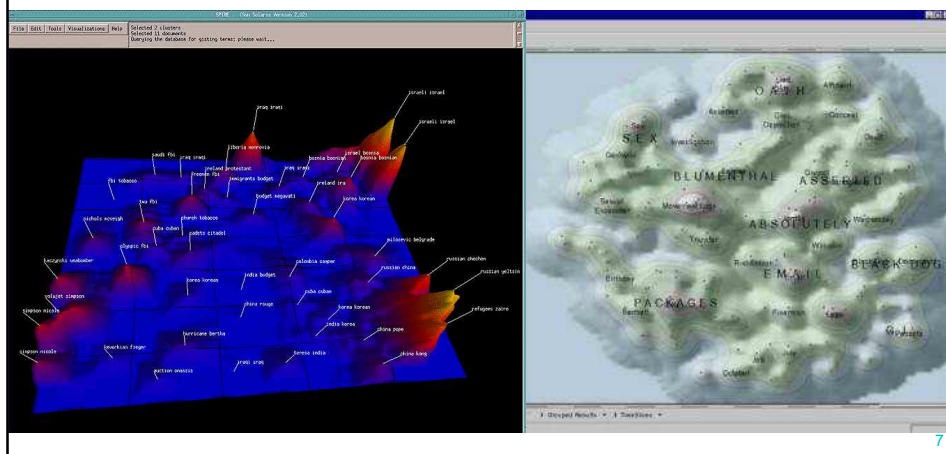
5

## Google News: automatic clustering gives an effective news presentation style



6

- » In this contour map, each mountain stands for a certain cluster (with its most frequent term) and mountain height stands for its cluster size

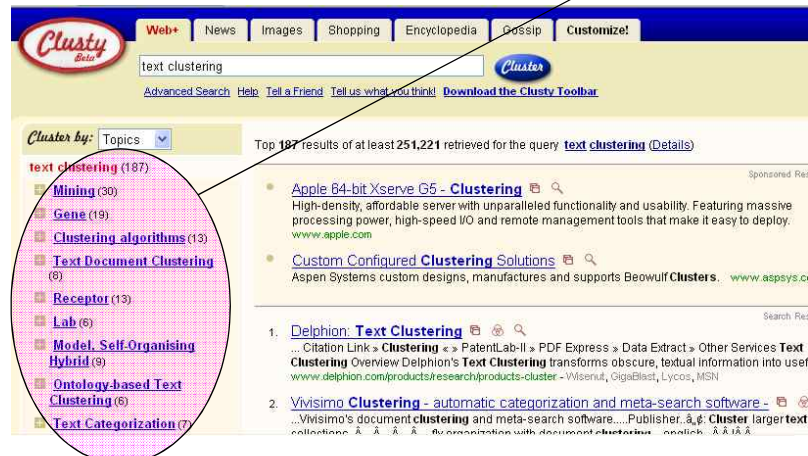


- **Basic Cluster hypothesis** : Documents in the same cluster behave **similarly** with respect to relevance to information needs
- Therefore, to improve recall in IR system
  - » Prior process: we can cluster all the documents in collection
  - » When a query matches a document  $D$  (that is,  $D$  is relevant to that query), IR system return other documents in the cluster containing that document  $D$
- By doing this, we hope that querying with “car” will also return documents containing “*automobile*”
  - » Because clustering groups documents containing “car” together with documents containing “*automobile*”

A document containing “*car*” is likely to have many common terms(or words) with a document containing “*automobile*” so that the similarity between the two documents is apt to be very high in vector space model

Clustering can be also used for better navigation of search results in IR system

- As shown in the picture, this search engine gives well-categorized search results that are grouped thematically by clustering



## Document clustering

- Given a set of documents, document clustering groups those documents into a number of clusters based on their content similarity

## Issues for document clustering

- Representation and similarity Metric for clustering
  - » Document representation
    - Is it based on vector space model or boolean model?
  - » Similarity or distance metrics
    - Is cosine similarity, Euclidean distance or anything else used ?
- How many clusters?
  - » Is it fixed a priori?
  - » Is it completely driven from document collection?
  - » Basic principle: just **avoid “trivial” cluster that is too large or too small**
    - In an application, if a cluster is too large or too small, then we generally **waste an extra user click without narrowing down the set of documents enough** when navigating document collection

11

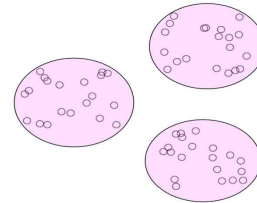
## What makes documents “related”?

- Ideally, semantic similarity makes documents related each other
- But, practically, **statistical similarity** makes documents related each other
  - » We can use **cosine similarity**
  - » To make it effective, documents are represented as **term vectors**
  - » Incidentally, for many IR algorithms, it is easier to think similarity in terms of a **distance** (rather than similarity) between documents
  - » We often use **Euclidean distance** as a similarity metric (more exactly, dissimilarity metric)


12

## Clustering Algorithms

- Flat algorithm
  - » Usually, it start with a random partitioning
  - » Gradually, it is refined iteratively
    - e.g. K-means clustering



- Hierarchical algorithms
  - » Bottom-up, agglomerative
  - » Top-down, divisive



13

## K-Means (flat clustering algorithm)

- It assumes that documents are real-valued term vectors
- Cluster is based on centroid (the center of gravity or mean) of term vectors:

In a cluster  $c$ ,  $\vec{\mu}(c) = \frac{1}{|c|} \sum_{\vec{x} \in c} \vec{x}$

- Re-assigning instances to clusters is based on Euclidean distance to each cluster centroid.

Note that henceforth, we can equivalently phrase Euclidean distance in terms of dissimilarity)

14

## K-Means Algorithm for making $K$ clusters

```


select  $K$  random docs  $\{s_1, s_2, \dots, s_K\}$  as seeds.
do
{
  for each doc  $d_i$ 
  {
    assign  $d_i$  to the cluster  $c_j$  such that  $dist(d_i, s_j)$  is minimal ;
  }
  for each cluster  $c_j$  // update the seeds to the centroid of each cluster
  {
     $s_j = \mu(c_j)$  ;
  }
} until (clustering is converged enough or termination condition is met)

```

Much simpler and easier  
than expected !

15

## K Means Example ( $K=2$ )



Pick seeds  
 Reassign clusters  
 Compute centroids  
 Reassign clusters  
 Compute centroids  
 Reassign clusters  
**Clustering Converged!**


16



## Termination conditions

- We may have several options for termination condition of K-means algorithm
- e.g.
  - » just terminate in a fixed number of iterations
  - » terminate when all clusters are unchanged at all
  - » terminate when all cluster centroid positions are unchanged at all

Does the second termination condition is equivalent to the third one?  
yes or no ?



17

## Convergence

- Why should the *K*-means algorithm ever reach a *fixed point*?
  - » There is always a state in which all clusters don't change
- Indirect Answer: *K*-means algorithm is a special case of a general procedure known as *Expectation Maximization (EM) algorithm*
  - » EM is proved to converge to a stable state  
(We assume you've seen EM algorithm in AI class before)
  - » Number of iterations for convergence can be large, but in practice it is usually not be large (much smaller than expected)

18

## Time Complexity of K-Means Algorithm

```

select  $K$  random docs  $\{s_1, s_2, \dots, s_K\}$  as seeds.
do
{
  for each doc  $d_i$ 
    assign  $d_i$  to the cluster  $c_j$  such that  $dist(d_i, s_j)$  is minimal ;
  for each cluster  $c_j$ 
     $s_j = \mu(c_j)$  ;
} until (clustering is converged enough)
    
```

- Computing distance between two term vectors is  $O(m)$  where  $m$  is the dimensionality of the vectors
- Reassigning clusters needs  $O(Kn)$  distance computations and thus  $O(Knm)$  where  $n$  is the number of all document vectors
- Computing centroids requires each document vector to be added once to some centroid and thus needs  $O(nm)$
- Because these two steps (reassigning clusters and computing centroids) are each performed once for  $I$  iterations, **the total time complexity is  $O(IKnm)$  → Totally linear complexity !**

19

## Initial Seed Choice in K-Means Algorithm

- Results can vary based on initial random seed selection.
- It means that some initial seeds can result in poor convergence rate, or convergence to a sub-optimal clustering state

### » Solutions

- Select good initial seeds using a heuristic (e.g. select a document vector that is most distant from any (already) selected seed)
- Try multiple times with different initial seeds because this algorithm apparently uses hill-climbing searching  
(We assume you've seen hill-climbing searching in algorithm or AI classes before)

### an example showing sensitivity to initial seeds

A	B	C
○	○	○
○	○	○
D	E	F

### In the above,

- if you start with B and E as centroids, you converge to {A,B,C} and {D,E,F}
- If you start with D and F, you converge to {A,B,D,E} and {C,F}


20

# Hierarchical Clustering Algorithm

- Two ways

- » Divisive Clustering (top down)

- The complete document collection is assumed to represent one complete cluster
- Then, each cluster is subsequently and recursively broken down into smaller clusters using flat clustering algorithm like *k*-means algorithm



- » Hierarchical Agglomerative Clustering (bottom up)

- Individual document similarities are used as a starting point.
- Based on it, a gluing operation agglomeratively collects similar documents (or similar document clusters) into a super-ordinate (or larger) cluster

21

## Hierarchical Agglomerative Clustering

- Basic procedure

1. Place each of  $N$  document vectors into a cluster of its own ;
2. Compute all pairwise cluster-to-cluster similarities ;
3. Form a new larger cluster by combining the pair of most similar clusters ;
4. Goto step 2 while the number of clusters left is greater than 1 ;

22

## How to Combine Clusters?

- **Basic procedure**

1. Place each of  $N$  document vectors into a cluster of its own ;
2. Compute all pairwise cluster-to-cluster similarities ;
3. Form a new larger cluster by combining the pair of most similar clusters ;
4. Goto step 2 if the number of clusters left is greater than 1 ;

- Just combine two clusters using intercluster similarity

- » Three options
  - Single-link clustering
  - Complete-link clustering
  - (Group) average-link clustering

23

## How to Combine Clusters?

- **Basic procedure**

1. Place each of  $N$  document vectors into a cluster of its own ;
2. Compute all pairwise cluster-to-cluster similarities ;
3. Form a new larger cluster by combining the pair of most similar clusters ;
4. Goto step 2 if the number of clusters left is greater than 1 ;

- Single-link clustering is based on the following three criteria

- » Each document vector must have a similarity exceeding a threshold value with at least one other document vector in the same cluster
- » The similarity between two clusters is taken to be the similarity between the most similar pair of document vectors in the two clusters
- » Each cluster member(=document vector) will be more similar to at least one member in that same cluster than to any member of another cluster

24

## How to Combine Clusters? *(Continued)*

- **Basic procedure**

1. Place each of  $N$  document vectors into a cluster of its own ;
2. Compute all pairwise cluster-to-cluster similarities ;
3. Form a new larger cluster by combining the pair of most similar clusters ;
4. Goto step 2 if the number of clusters left is greater than 1 ;

- **Complete-link clustering is based on the following three criteria**
  - » Each document vector must have a similarity exceeding a threshold value with **all other document vectors** in the same cluster
  - » The similarity between two clusters is taken to be the similarity between **the least similar pair** of document vectors in the two clusters
  - » Each cluster member(=document vector) is more similar to **the most dissimilar member** of that same cluster than to the **most dissimilar member** of any other cluster

25

## How to Combine Clusters? *(Continued)*

- **Basic procedure**

1. Place each of  $N$  document vectors into a cluster of its own ;
2. Compute all pairwise cluster-to-cluster similarities ;
3. Form a new larger cluster by combining the pair of most similar clusters ;
4. Goto step 2 if the number of clusters left is greater than 1 ;

- **Group average-link clustering**
  - » Makes a compromise between the extreme stuffs of single-link and complete-link clusterings
  - » Each cluster member(=document vector) has a **greater average similarity** to the remaining members of that same cluster than it does to all members of any other cluster

26

## Single-link, complete-link, and group average-link clusterings

- Note: for easy understanding of single-link, complete-link, and group average-link clusterings, we may need some examples and try out each of them for ourselves

27

## Detailed examples for single-link clusterings in terms of run-time data structure

A-F (6 items)  $6(6-1)/2$  (15) pairwise similarities

Rank	Pair	Similarity	Rank	Pair	Similarity
1	AF	0.9	9	BC	0.4
2	AE	0.8	10	DE	0.4
3	BF	0.8	11	AB	0.3
4	BE	0.7	12	CD	0.3
5	AD	0.6	13	EF	0.3
6	AC	0.5	14	CF	0.2
7	BD	0.5	15	DF	0.1
8	CE	0.5			

decreasing order

Actually only a upper or lower triangular matrix is enough !




	A	B	C	D	E	F
A	.	.3	.5	.6	.8	.9
B	.3	.	.4	.5	.7	.8
C	.5	.4	.	.3	.5	.2
D	.6	.5	.3	.	.4	.1
E	.8	.7	.5	.4	.	.3
F	.9	.8	.2	.1	.3	.

initial data structure

28

## Single Link Clustering

1. AF: 0.9




$$\text{sim}(AF, X) = \max(\text{sim}(A, X), \text{sim}(F, X))$$

	A	B	C	D	E	F
A	.	.3	.5	.6	.8	.9
B	.3	.	.4	.5	.7	.8
C	.5	.4	.	.3	.5	.2
D	.6	.5	.3	.	.4	.1
E	.8	.7	.5	.4	.	.3
F	.9	.8	.2	.1	.3	.

largest one

2. AF, E: 0.8



$$\text{sim}(AEF, X) = \max(\text{sim}(AF, X), \text{sim}(E, X))$$

	AE	B	C	D	E
AE	.	.8	.5	.6	.8
B	.8	.	.4	.5	.7
C	.5	.4	.	.3	.5
D	.6	.5	.3	.	.4
E	.8	.7	.5	.4	.


largest one

29

## Single Link Clustering (Cont.)

3. AEF, B: 0.8

Note E and B are on the same level.




$$\text{sim}(ABEF, X) = \max(\text{sim}(AEF, X), \text{sim}(B, X))$$

	AEF	B	C	D
AEF	.	.8	.5	.6
B	.8	.	.4	.5
C	.5	.4	.	.3
D	.6	.5	.3	.

largest one

4. ABDEF, D: 0.6



$$\text{sim}(ABDEF, X) = \max(\text{sim}(ABEF, X), \text{sim}(D, X))$$


	ABEF	C	D
ABEF	.	.5	.6
C	.5	.	.3
D	.6	.3	.

largest one

30

## Single Link Clustering (Cont.)

5. ABDEF, C: 0.5



31

## Single-Link Clusters

- Similarity level 0.85 (i.e., similarity threshold)

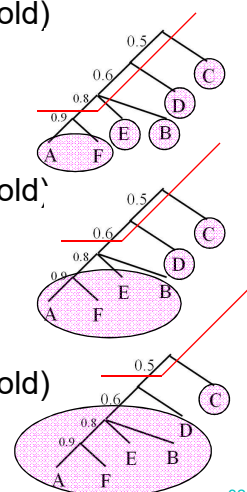
» AF, E, B, D, C

- Similarity level 0.65 (i.e., similarity threshold)

» AFEB, D, C

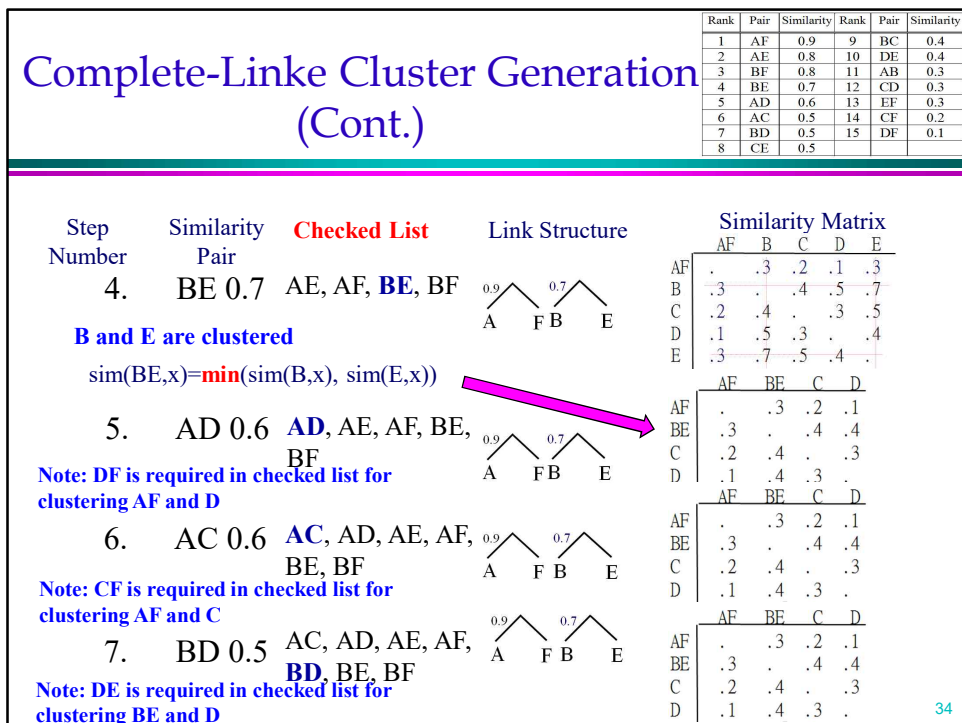
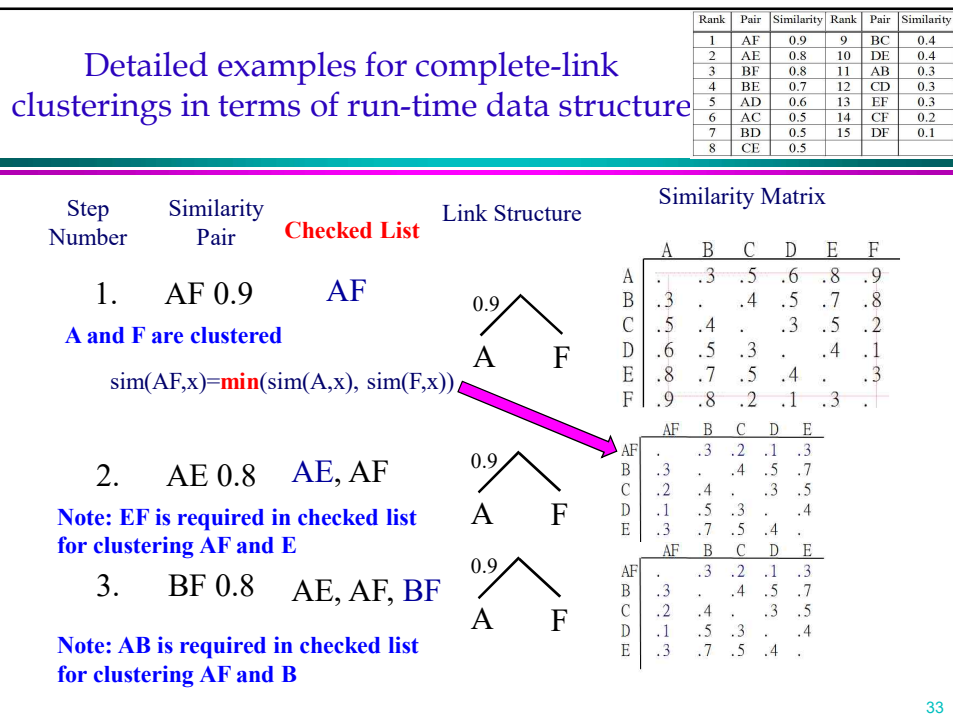
- Similarity level 0.55 (i.e., similarity threshold)

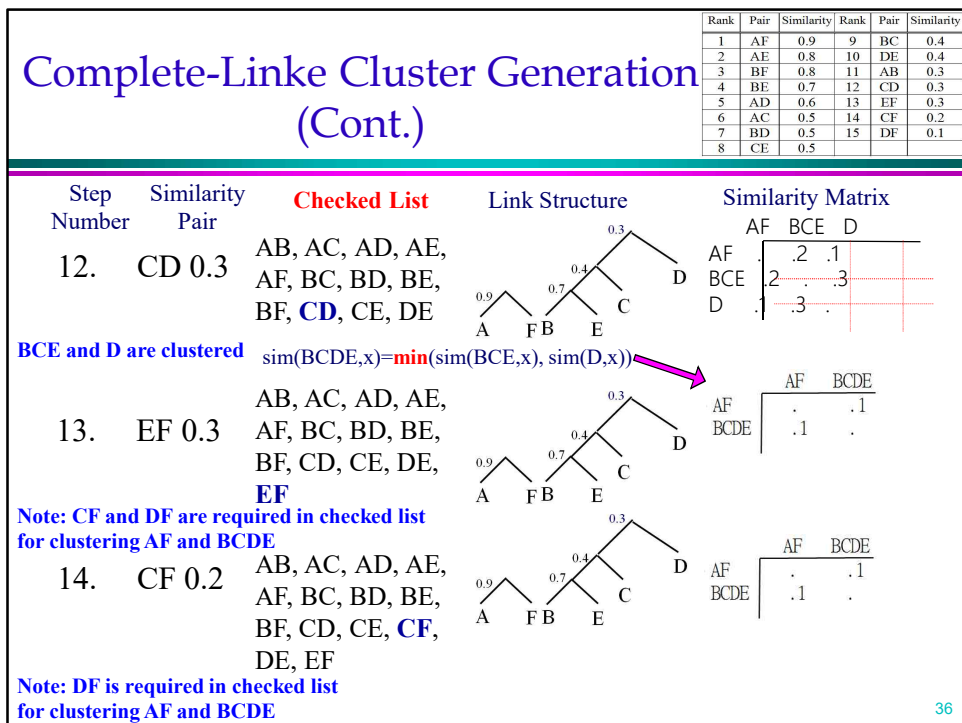
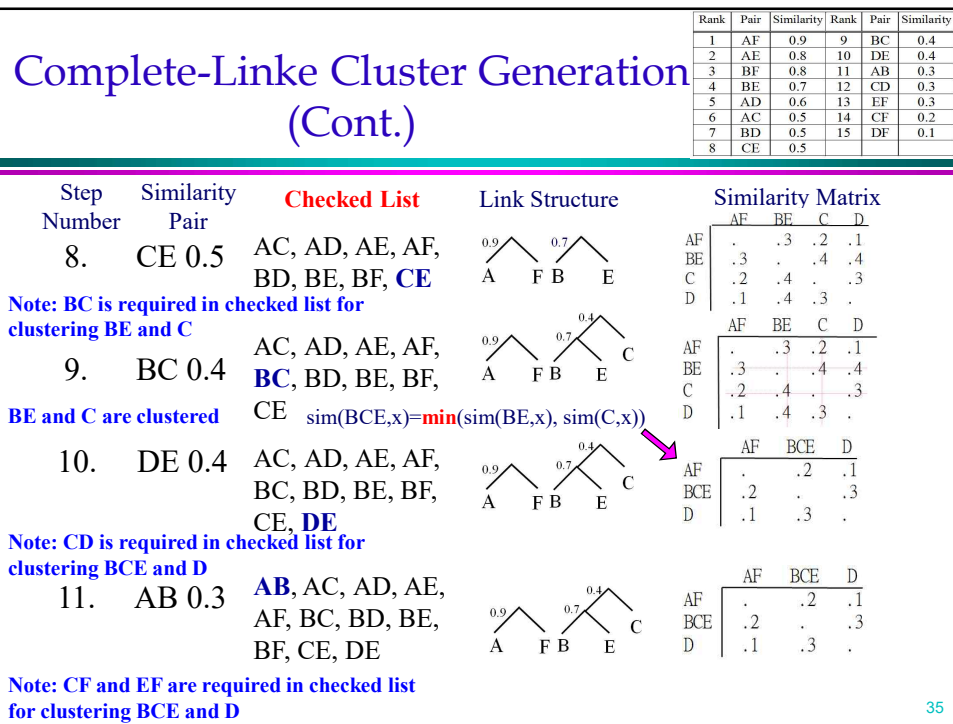
» AFEBD, C



32







## Complete-Link Cluster Generation (Cont.)

Rank	Pair	Similarity	Rank	Pair	Similarity
1	AF	0.9	9	BC	0.4
2	AE	0.8	10	DE	0.4
3	BF	0.8	11	AB	0.3
4	BE	0.7	12	CD	0.3
5	AD	0.6	13	EF	0.3
6	AC	0.5	14	CF	0.2
7	BD	0.5	15	DF	0.1
8	CE	0.5			


Step Number

15. DF 0.1

Checked List

AB, AC, AD, AE,  
AF, BC, BD, BE,  
BF, CD, CE, CF,  
DE, **DF**, EF

Link Structure



Similarity Matrix

AF	BCDE
AF	.1
BCDE	.1

BCDE and AF are clustered

$$\text{sim}(\text{ABCDEF}, x) = \min(\text{sim}(\text{BCDE}, x), \text{sim}(\text{AF}, x))$$

Actually, we don't need to do the final update of matrix

37

## Complete Link Clusters

- Similarity level 0.85 (i.e., similarity threshold)




» AF, B, E, C, D

- Similarity level 0.65 (i.e., similarity threshold)

» AF, BE, C, D

- Similarity level 0.35 (i.e., similarity threshold)

» AF, BCE, D



38

## Group Average Link Clustering

- This lecture will not give a detailed example of group average-link clustering
  - » But, in order to formulate the details of such clustering with their data structure
    - Just, use the average values of pair-wise members within a cluster to update the similarity matrix
    - remember that all cluster members contribute to inter-cluster similarity
    - It usually result in a **intermediately bound structure** between the **loosely bound** single-link clusters and **tightly bound** complete-link clusters
  - » Formulate more details of group average-link clustering by yourself, if necessary

39

## Comparison

- The Behavior of Single-Link Clustering
  - » The single-link clustering process tends to produce **a small number of large and loosely linked clusters** that are characterized by a chaining effect → This often results in unfair growth of clusters
  - » Each element is usually attached to only one other member of the same cluster at each similarity level
  - » It is sufficient to remember the step sequence of previously clustered example when using single-link clustering

40

## Comparison

- The Behavior of Complete-Link Clustering
  - » Complete-link clustering process tends to produce a much larger number of small and tightly linked groupings
  - » It is because each member in a complete-link cluster is firmly guaranteed to resemble all other members in that cluster at the stated similarity level
  - » It is necessary to remember the step sequence of previously clustered example when using complete-link clustering
- Comparison
  - » Generally, the complete-link clustering may be better adapted to information retrieval than the single-link clustering
  - » Of course, a complete-link cluster generation is more expensive to perform than a comparable single-link process